

Hacks for training RL systems from John Schulman's lecture at Deep RL Bootcamp (Aug 2017)

 8 commits 1 branch 0 releases 1 contributor

Branch: master ▾


New pull request

Create new file

Upload files

Find file

Clone or download ▾

 **williamFalcon** committed on **GitHub** Update README.md

Latest commit c2b81e7 11 hours ago

 [README.md](#)

Update README.md

11 hours ago

 README.md

# DeepRLHacks

From a talk given by [John Schulman](#) titled "The Nuts and Bolts of Deep RL Research" (Aug 2017)

These are tricks written down while attending summer [Deep RL Bootcamp at UC Berkeley](#).

## Tips to debug new algorithm

1. Simplify the problem by using a low dimensional state space environment.
  - John suggested to use the [Pendulum problem](#) because the problem has a 2-D state space (angle of pendulum and velocity).
  - Easy to visualize what the value function looks like and what state the algorithm should be in and how they evolve over time.
  - Easy to visually spot why something isn't working (aka, is the value function smooth enough and so on).
2. To test if your algorithm is reasonable, construct a problem you know it should work on.
  - Ex: For hierarchical reinforcement learning you'd construct a problem with an OBVIOUS hierarchy it should learn.
  - Can easily see if it's doing the right thing.
  - WARNING: Don't over fit method to your toy problem (realize it's a toy problem).
3. Familiarize yourself with certain environments you know well.
  - Over time, you'll learn how long the training should take.
  - Know how rewards evolve, etc...
  - Allows you to set a benchmark to see how well you're doing against your past trials.
  - John uses the hopper robot where he knows how fast learning should take, and he can easily spot odd behaviors.

## Tips to debug a new task

1. Simplify the task
  - Start simple until you see signs of life.
  - Approach 1: Simplify the feature space:
    - For example, if you're learning from images (huge dimensional space), then maybe hand engineer features first. Example: If you think your function is trying to approximate a location of something, use the x,y location as features as step 1.
    - Once it starts working, make the problem harder until you solve the full problem.
  - Approach 2: simplify the reward function.

- Formulate so it can give you FAST feedback to know whether you're doing the right thing or not.
- Ex: Have reward for robot when it hits the target (+1). Hard to learn because maybe too much happens in between starting and reward. Reformulate as distance to target instead which will increase learning and allow you to iterate faster.

## Tips to frame a problem in RL

---

Maybe it's unclear what the features are and what the reward should be, or if it's feasible at all.

1. First step: Visualize a random policy acting on this problem.
  - See where it takes you.
  - If random policy on occasion does the right thing, then high chance RL will do the right thing.
    - Policy gradient will find this behavior and make it more likely.
  - If random policy never does the right thing, RL will likely also not.
2. Make sure observations usable:
  - See if YOU could control the system by using the same observations you give the agent.
    - Example: Look at preprocessed images yourself to make sure you don't remove necessary details or hinder the algorithm in a certain way.
3. Make sure everything is reasonably scaled.
  - Rule of thumb:
    - Observations: Make everything mean 0, standard deviation 1.
    - Reward: If you control it, then scale it to a reasonable value.
      - Do it across ALL your data so far.
  - Look at all observations and rewards and make sure there aren't crazy outliers.
4. Have good baseline whenever you see a new problem.
  - It's unclear which algorithm will work, so have a set of baselines (from other methods)
    - Cross entropy method
    - Policy gradient methods
    - Some kind of Q-learning method (checkout [OpenAI Baselines](#) as a starter or [RLLab](#))

## Reproducing papers

---

Sometimes (often), it's hard to reproduce results from papers. Some tricks to do that:

1. Use more samples than needed.
2. Policy right... but not exactly
  - Try to make it work a little bit.
  - Then tweak hyper parameters to get up to the public performance.
  - If want to get it to work at ALL, use bigger batch sizes.
    - If batch size is too small, noisy will overpower signal.
    - Example: TRPO, John was using too tiny of a batch size and had to use 100k time steps.
    - For DQN, best hyperparams: 10k time steps, 1mm frames in replay buffer.

## Guidelines on-going training process

---

Sanity check that your training is going well.

1. Look at sensitivity of EVERY hyper parameter

- If algo is too sensitive, then NOT robust and should NOT be happy with it.
- Sometimes it happens that a method works one way because of funny dynamics but NOT in general.

## 2. Look for indicators that the optimization process is healthy.

- Varies
- Look at whether value function is accurate.
  - Is it predicting well?
  - Is it predicting returns well?
  - How big are the updates?
- Standard diagnostics from deep networks

## 3. Have a system for continuously benchmarking code.

- Needs DISCIPLINE.
- Look at performance across ALL previous problems you tried.
  - Sometimes it'll start working on one problem but mess up performance in others.
  - Easy to over fit on a single problem.
- Have a battery of benchmarks you run occasionally.

## 4. Think your algorithm is working but you're actually seeing random noise.

- Example: Graph of 7 tasks with 3 algorithms and looks like 1 algorithm might be doing best on all problems, but turns out they're all the same algorithm with DIFFERENT random seeds.

## 5. Try different random seeds!!

- Run multiple times and average.
- Run multiple tasks on multiple seeds.
  - If not, you're likely to over fit.

## 6. Additional algorithm modifications might be unnecessary.

- Most tricks are ACTUALLY normalizing something in some way or improving your optimization.
- A lot of tricks also have the same effect... So you can remove some of them and SIMPLIFY your algorithm (VERY KEY).

## 7. Simplify your algorithm

- Will generalize better

## 8. Automate your experiments

- Don't spend your whole day watching your code spit out numbers.
- Launch experiments on cloud services and analyze results.
- Frameworks to track experiments and results:
  - Mostly uses iPython notebooks.
  - DBs seem unnecessary to store results.

# General training strategies

## 1. Whiten and standardize data (for ALL seen data since the beginning).

- Observations:
  - Do it by computing a running mean and standard deviation. Then z-transform everything.
  - Over ALL data seen (not just the recent data).
    - At least it'll scale down over time how fast it's changing.

- Might trip up the optimizer if you keep changing the objective.
- Rescaling (by using recent data) means your optimizer probably didn't know about that and performance will collapse.
- Rewards:
  - Scale and DON'T shift.
    - Affects agent's will to live.
    - Will change the problem (aka, how long you want it to survive).
- Standardize targets:
  - Same way as rewards.
- PCA Whitening?
  - Could help.
  - Starting to see if it actually helps with neural nets.
  - Huge scales (-1000, 1000) or (-0.001, 0.001) certainly makes learning slow.

## 2. Parameters that inform discount factors.

- Determines how far you're giving credit assignment.
- Ex: if factor is 0.99, then you're ignoring what happened 100 steps ago... Means you're shortsighted.
  - Better to look at how that corresponds to real time
    - Intuition, in RL we're usually discretizing time.
    - aka: are those 100 steps 3 seconds of actual time?
    - what happens during that time?
- If TD methods for policy gradient of Value fx estimation, gamma can be close to 1 (like 0.999)
  - Algo becomes very stable.

## 3. Look to see that problem can actually be solved in the discretized level.

- Example: In game if you're doing frame skip.
  - As a human, can you control it or is it impossible?
  - Look at what random exploration looks like
    - Discretization determines how far your browsing motion goes.
    - If do many actions in a row, then tend to explore further.
    - Choose your time discretization in a way that works.

## 4. Look at episode returns closely.

- Not just mean, look at min and max.
  - The max return is something your policy can hone in pretty well.
  - Is your policy ever doing the right thing??
- Look at episode length (sometimes more informative than episode reward).
  - if on game you might be losing every time so you might never win, but... episode length can tell you if you're losing SLOWER.
  - Might see a episode length improvement in the beginning but maybe not reward.

# Policy gradient diagnostics

## 1. Look at entropy really carefully

- Entropy in ACTION space
  - Care more about entropy in state space, but don't have good methods for calculating that.
- If going down too fast, then policy becoming deterministic and will not explore.

- If NOT going down, then policy won't be good because it is really random.
- Can fix by:
  - KL penalty
    - Keep entropy from decreasing too quickly.
  - Add entropy bonus.
- How to measure entropy.
  - For most policies can compute entropy analytically.
    - If continuous usually using a Gaussian so can compute differential entropy.

## 2. Look at KL divergence

- Look at size of updates in terms of KL divergence.
- example:
  - If KL is .01 then very small.
  - If 10 then too much.

## 3. Baseline explained variance.

- See if value function is actually a good predictor or a reward.
  - if negative it might be over fitting or noisy.
    - Likely need to tune hyper parameters

## 4. Initialize policy

- Very important (more so than in supervised learning).
- Zero or tiny final layer to maximize entropy
  - Maximize random exploration in the beginning

## Q-Learning Strategies

---

1. Be careful about replay buffer memory usage.
  - You might need a huge buffer, so adapt code accordingly.
2. Play with learning rate schedule.
3. If converges slowly or has slow warm-up period in the beginning
  - Be patient... DQN converges VERY slowly.

## Bonus from [Andrej Karpathy](#):

---

1. A good feature can be to take the difference between two frames.
  - This delta vector can highlight slight state changes otherwise difficult to distinguish.